

Algorithms Implemented

- Monte Carlo Control (Exploring Starts initially)
 - Monte Carlo Control (ϵ -greedy / Epsilon-Soft version)
 - SARSA (On-Policy TD Learning)
 - Q-Learning (Off-Policy TD Learning)
-
-

Results

Algorithm	Average Return	Success Rate (%)	Final Policy (Best Actions)
MC-ES (Exploring Starts)	0.00	0%	All 0s (untrained)
Q-Learning	0.8	78-95%	[[2 2 1 0] [1 0 1 0] [2 1 1 0] [0 2 2 0]]
SARSA	0.7111	71.11-90%	[[1 2 1 0] [1 0 1 0] [2 2 1 0] [0 2 2 0]]
MC-ES (Epsilon-Soft)	0.78	78-90%	

: Different runs and
varying gamma
(discount factor) values
led to different
outcomes here are
some:

```
MC-ES -> Avg Success Rate: 0.00% ± 0.00%
MC-ES Policy (Last Run)
[[0 0 0 0]
 [0 0 0 0]
 [0 0 0 0]
 [0 0 0 0]]

Q-Learning -> Avg Success Rate: 100.00% ± 0.00%
Q-Learning Policy (Last Run)
[[2 2 1 0]
 [3 0 1 0]
 [2 2 1 0]
 [0 2 2 0]]

SARSA -> Avg Success Rate: 100.00% ± 0.00%
SARSA Policy (Last Run)
[[1 2 1 0]
 [1 0 1 0]
 [2 2 1 0]
 [0 2 2 0]]
```

```
(base) tazmeen@afroz:~$ python3 rl2.py
```

```
--- Run 1/1 ---
```

```
Training with Q-Learning...
```

```
Q-Learning Episode 0/10000
```

```
Q-Learning Episode 1000/10000
```

```
Q-Learning Episode 2000/10000
```

```
Q-Learning Episode 3000/10000
```

```
Q-Learning Episode 4000/10000
```

```
Q-Learning Episode 5000/10000
```

```
Q-Learning Episode 6000/10000
```

```
Q-Learning Episode 7000/10000
```

```
Q-Learning Episode 8000/10000
```

```
Q-Learning Episode 9000/10000
```

```
Q-Learning Average Return: 0.7111
```

```
Success rate: 71.11%
```

```
Training with SARSA...
```

```
SARSA Episode 0/10000
```

```
SARSA Episode 1000/10000
```

```
SARSA Episode 2000/10000
```

```
SARSA Episode 3000/10000
```

```
SARSA Episode 4000/10000
```

```
SARSA Episode 5000/10000
```

```
SARSA Episode 6000/10000
```

```
SARSA Episode 7000/10000
```

```
SARSA Episode 8000/10000
```

```
SARSA Episode 9000/10000
```

```
SARSA Average Return: 0.7111
```

```
Success rate: 71.11%
```

```
Training with Monte Carlo ES...
```

```
MC-ES Episode 0/10000
```

```
MC-ES Episode 1000/10000
```

```
MC-ES Episode 2000/10000
```

```
MC-ES Episode 3000/10000
```

```
MC-ES Episode 4000/10000
```

```
MC-ES Episode 5000/10000
```

```
MC-ES Episode 6000/10000
```

```
MC-ES Episode 7000/10000
```

```
MC-ES Episode 8000/10000
```

```
MC-ES Episode 9000/10000
```

```
MC-ES Final Success Rate: 98.50%
```

```
MC-ES -> Avg Success Rate: 100.00% ± 0.00%
```

```
MC-ES Policy (Last Run)
```

```
[[2 2 1 0]
```

```
 [1 0 1 0]
```

```
 [2 1 1 0]
```

```
 [0 2 2 0]]
```

```
SARSA Episode 8000/10000
SARSA Episode 9000/10000
SARSA Average Return: 0.7111
Success rate: 71.11%
```

```
Training with Monte Carlo ES...
```

```
MC-ES Episode 0/10000
MC-ES Episode 1000/10000
MC-ES Episode 2000/10000
MC-ES Episode 3000/10000
MC-ES Episode 4000/10000
MC-ES Episode 5000/10000
MC-ES Episode 6000/10000
MC-ES Episode 7000/10000
MC-ES Episode 8000/10000
MC-ES Episode 9000/10000
MC-ES Final Success Rate: 98.50%
```

```
MC-ES -> Avg Success Rate: 100.00% ± 0.00%
```

```
MC-ES Policy (Last Run)
```

```
[[2 2 1 0]
 [1 0 1 0]
 [2 1 1 0]
 [0 2 2 0]]
```

```
Q-Learning -> Avg Success Rate: 100.00% ± 0.00%
```

```
Q-Learning Policy (Last Run)
```

```
[[2 2 1 0]
 [1 0 1 0]
 [2 2 1 0]
 [0 2 2 0]]
```

```
SARSA -> Avg Success Rate: 100.00% ± 0.00%
```

```
SARSA Policy (Last Run)
```

```
[[2 2 1 0]
 [1 0 1 0]
 [2 2 1 0]
 [0 2 2 0]]
```

```
Creating animation for MC-ES...
```

```
Success rate over 10 episodes: 100.0%
```

```
Animation saved to MC-ES_animation.gif
```

```
Creating animation for Q-Learning...
```

```
Success rate over 10 episodes: 100.0%
```

```
Animation saved to Q-Learning_animation.gif
```

```
Creating animation for SARSA...
```

```
Success rate over 10 episodes: 100.0%
```

```
Animation saved to SARSA_animation.gif
```



```
===== RESULTS SUMMARY =====  
Number of episodes: 10000  
Environment: FrozenLake-v1 (is_slippery=False)
```

```
Monte Carlo ES:  
  Average Return: 0.0000  
  Success Rate: 0.00%  
  Total Time: 89.82 seconds
```

```
Monte Carlo  $\epsilon$ -Soft:  
  Average Return: 0.3529  
  Success Rate: 35.29%  
  Total Time: 58.36 seconds
```

```
Q-Learning:  
  Average Return: 0.6989  
  Success Rate: 69.89%  
  Total Time: 5.28 seconds
```

```
SARSA:  
  Average Return: 0.0000  
  Success Rate: 0.00%  
  Total Time: 15.45 seconds
```

```
(base) tazmeen@afroz:~$
```

```
===== RESULTS SUMMARY =====  
Number of episodes: 10000  
Environment: FrozenLake-v1 (is_slippery=False)
```

```
Monte Carlo ES:  
  Average Return: 0.0000  
  Success Rate: 0.00%  
  Total Time: 74.12 seconds
```

```
Monte Carlo  $\epsilon$ -Soft:  
  Average Return: 0.4298  
  Success Rate: 42.98%  
  Total Time: 44.52 seconds
```

```
Q-Learning:  
  Average Return: 0.0962  
  Success Rate: 9.62%  
  Total Time: 12.34 seconds
```

```
SARSA:  
  Average Return: 0.0000  
  Success Rate: 0.00%  
  Total Time: 16.68 seconds
```

```
(base) tazmeen@afroz:~$
```

```

===== RESULTS SUMMARY =====
Number of episodes: 10000
Environment: FrozenLake-v1 (is_slippery=False)

Monte Carlo ES:
  Average Return: 0.0000
  Success Rate: 0.00%
  Total Time: 60.24 seconds

Monte Carlo  $\epsilon$ -Soft:
  Average Return: 0.7829
  Success Rate: 78.29%
  Total Time: 29.53 seconds

Q-Learning:
  Average Return: 0.0000
  Success Rate: 0.00%
  Total Time: 11.55 seconds

SARSA:
  Average Return: 0.7587
  Success Rate: 75.87%
  Total Time: 2.15 seconds

(base) tazmeen@afroz:~$ gedit test8.py
(base) tazmeen@afroz:~$ gedit test9.py
(base) tazmeen@afroz:~$ gedit test10.py

```

(The original script did not include animation code. Animation was generated separately using code provided by GPT)

Function to create and save animation

```
def create_animation(frames, filename="agent_animation.gif", fps=4):
```

```
    fig, ax = plt.subplots(figsize=(5, 5))
    plt.tight_layout()
    plt.axis('off')
```

Create animation

```
    patch = ax.imshow(frames[0])
```

```
    def animate(i):
```

```
        patch.set_array(frames[i])
```

```
        return [patch]
```

```
    anim = animation.FuncAnimation(
```

```
        fig, animate, frames=len(frames), interval=1000//fps)
```

Save animation

```
    anim.save(filename, writer=PillowWriter(fps=fps))
```

```
    plt.close()
```

```
    print(f"Animation saved to {filename}")
```

```
    return filename
```